



## MASTER RESEARCH INTERNSHIP



## BIBLIOGRAPHIC REPORT

---

# Software Fault Isolation using the CompCert compiler

---

**Domaine: Cryptography and Security**

*Author:*  
Alexandre DANG

*Supervisor:*  
Frédéric BESSON  
Team CELTIQUE

**Abstract:**

## **Contents**

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b><i>Software Fault Isolation</i></b>	<b>1</b>
2.1	Principle . . . . .	1
2.1.1	Code generator . . . . .	2
2.1.2	Code verifier . . . . .	2
2.1.3	Pros and cons . . . . .	2
2.1.4	Implementations . . . . .	2
2.2	SFI using CompCert . . . . .	2
2.2.1	CompCert the verified compiler . . . . .	2
2.2.2	SFI with CompCert . . . . .	2
2.2.3	Evaluation of the approach . . . . .	3
2.3	Limits of SFI . . . . .	3
2.3.1	Return addresses . . . . .	3
2.3.2	Proposed solution . . . . .	3
<b>3</b>	<b>Protecting the stack with fixed frame size</b>	<b>3</b>
3.1	Description of the approach . . . . .	3
3.1.1	Issue . . . . .	3
3.1.2	Proposed solution . . . . .	3
3.2	Security properties . . . . .	6
3.3	Analysis of the approach . . . . .	7
3.3.1	Conditions . . . . .	7
3.3.2	Discussion . . . . .	8
<b>4</b>	<b>Implementation</b>	<b>8</b>
4.1	Implementation . . . . .	9
4.1.1	CompCert stack . . . . .	9
4.1.2	Fixed stack frames size . . . . .	10
4.1.3	Stack alignment . . . . .	11
4.1.4	Detection of memory write statements . . . . .	11
4.1.5	Securing memory write statements . . . . .	11
4.2	Evaluation of the implementation . . . . .	11

# 1 Introduction

- Secure malicious code through software solution
- Usage in applications which use modules from unknown origin (browsers, computer clusters)
- current appeal for SFI speed and small TCB
- SFI is still incomplete, especially with ROP attack => our approach
- plan

## 2 *Software Fault Isolation*

We introduce here *Software Fault Isolation* (SFI) which inspired us the idea to protect return addresses through fixed stack frame size. SFI aims to protect a main program from the different modules that he will need to use. These modules will be loaded in the same memory space as the main program but in a confined area called *sandbox*. The SFI mechanism is composed of two elements: a code generator and a verifier. The generator transforms the assembly code of the hazardous modules so that they will be constrained in the sandbox. The verifier operates just before loading the modules in the memory. It checks the if SFI transformations introduced by the generator are still present and valid. For the rest of the document we will reserve the word "program" to refer to the code protected by SFI and "module" to refer to the hazardous code.

### 2.1 Principle

The main principle behind SFI was first presented in the work of Wahbe and al. [ref](#). Later works [ref](#), which will be introduced in the chapter 2.1.4, are all based on the foundations of SFI detailed here. The implementation described here was realised for a RISC architecture like MIPS or *Alpha*.

SFI considers that a malicious code is effectively contained in the sandbox if these three security properties hold true:

- **Verified code**, only instructions that have been checked by the verifier will be executed
- **Memory safety**, malicious modules won't do any *write* or *jump* operations out of the sandbox
- **Flow control integrity**, every flow control transfer from hazardous modules to the main program is identified and verified

The first property protects us against self-modifying code which could bypass the SFI measures. *Memory safety* prevents any illegal access to the memory of the protected program. The last property allows us to authorized only licit interactions between the program and its modules. SFI forbids any call from malicious modules that could modify the flow control of the program. If the flow control was fiddled with, it could lead to an unexpected behaviour of the program which we want to avoid.

The code generator transforms the assembly code of the hazardous modules so that respect the security properties presented before. The generator is integrated to the compiler which will create

*sandboxed executable*. Afterwards this executable will be checked by the verifier before being loaded in the memory. It verifies that the transformations introduced by the generator are present and valid. If the verification fails the module will be rejected and won't be executed. We can note that we only need to trust the verifier to prevent running any dangerous module. It's one advantage of SFI, only the verifier needs to be in the *Trusted Computing base* (TCB).

### 2.1.1 Code generator

To protect the program from its modules, the generator will restrain every write and jump instructions of the modules to addresses of their sandbox. The generator has to face three issues to do so. Firstly, is to introduce protection mechanisms before every dangerous instructions. For example, assessing that the address of a jump instruction is an authorized one. Secondly, we have to make sure that these protection mechanisms can't be avoided. Finally, the transformations introduced have to authorized only legal calls out of the sandbox by using entry points specified by the protected program. For example, Google Chrome only allows its modules to use a specific interface to interact with the browser. This way the modules can't disrupt the flow control of Google Chrome easily.

**Confining memory accesses** The main program memory should avoid being corrupted by its modules. SFI aims to isolate these modules in a reserved of the program's memory called sandbox. The sandbox is a contiguous memory area which size is a power of two. Indeed, these requirements ease the confinement of the modules in their sandbox by using arithmetic operations on bits which accelerates the process.

## Protection of sandboxing mechanisms

### Controlled interactions with the protected program

#### 2.1.2 Code verifier

#### 2.1.3 Pros and cons

#### 2.1.4 Implementations

NativeClient, SFI for Google Chrome

## 2.2 SFI using CompCert

### 2.2.1 CompCert the verified compiler

CompCert

Memory model of CompCert

### 2.2.2 SFI with CompCert

Cminor

## Specification of the SFI transformation

### Masking in CompCert

#### 2.2.3 Evaluation of the approach

### 2.3 Limits of SFI

#### 2.3.1 Return addresses

#### 2.3.2 Proposed solution

## 3 Protecting the stack with fixed frame size

A lot of attacks on software aims to interfere with the control flow of the targeted program. Among those, *Returned Oriented Programming* (ROP) attacks specifically try to overwrite the return addresses. By doing so the attacked function will return to a malicious piece of code that will get executed (see Figure 1). Stack overflow is an example of such ROP attacks. We propose a solution against ROP attacks which combined with SFI would protect from most of control-flow interference attacks. Inspired from SFI techniques we aim to prevent any overwriting of the return addresses. To do so we need to know these return addresses location in the memory. Therefore our approach consists of modifying the stack structure in order to have an easy way to know the return addresses locations. With this knowledge we will be able to put a mask, as in SFI, before every dangerous write instructions and prevent any ROP attack.

### 3.1 Description of the approach

#### 3.1.1 Issue

We want to protect the return addresses in the stack from being overwritten illegally. The only moment they should be written over is during a function call routine. We want to use a masking operation similar as in SFI, therefore we need to be able to check if an address is the location of a return address. One solution to enable these runtime checks is to enforce a strict control over the stack, thus allowing us to know the exact locations of the return addresses. *suggest another idea maybe*

#### 3.1.2 Proposed solution

To be able to know the return addresses location we suggest a solution. The idea is to define a constant size for the frames of the stack. Stack frames are created and added to the stack after every function call, whereas they are removed from the stack when returning from the calling function. The frames are piled up on the stack following the FIFO rule (*First In First Out*). In these frames the return addresses usually have a specific place, and by fixing the frames size we impose the return addresses to have a fixed offset between them. With this property we can guess the location of every return address which would be:  $c \bmod n$  with  $c$  a constant and  $n$  the size of the frames. Now that we know  $n$  (the size of the frames), we need to get the constant  $c$ . To solve this, our idea is to align the beginning of the stack when starting the program. This way we can choose the constant  $c$  and we arrange it to have  $c = 0$ . More details can be seen on Figure 2.

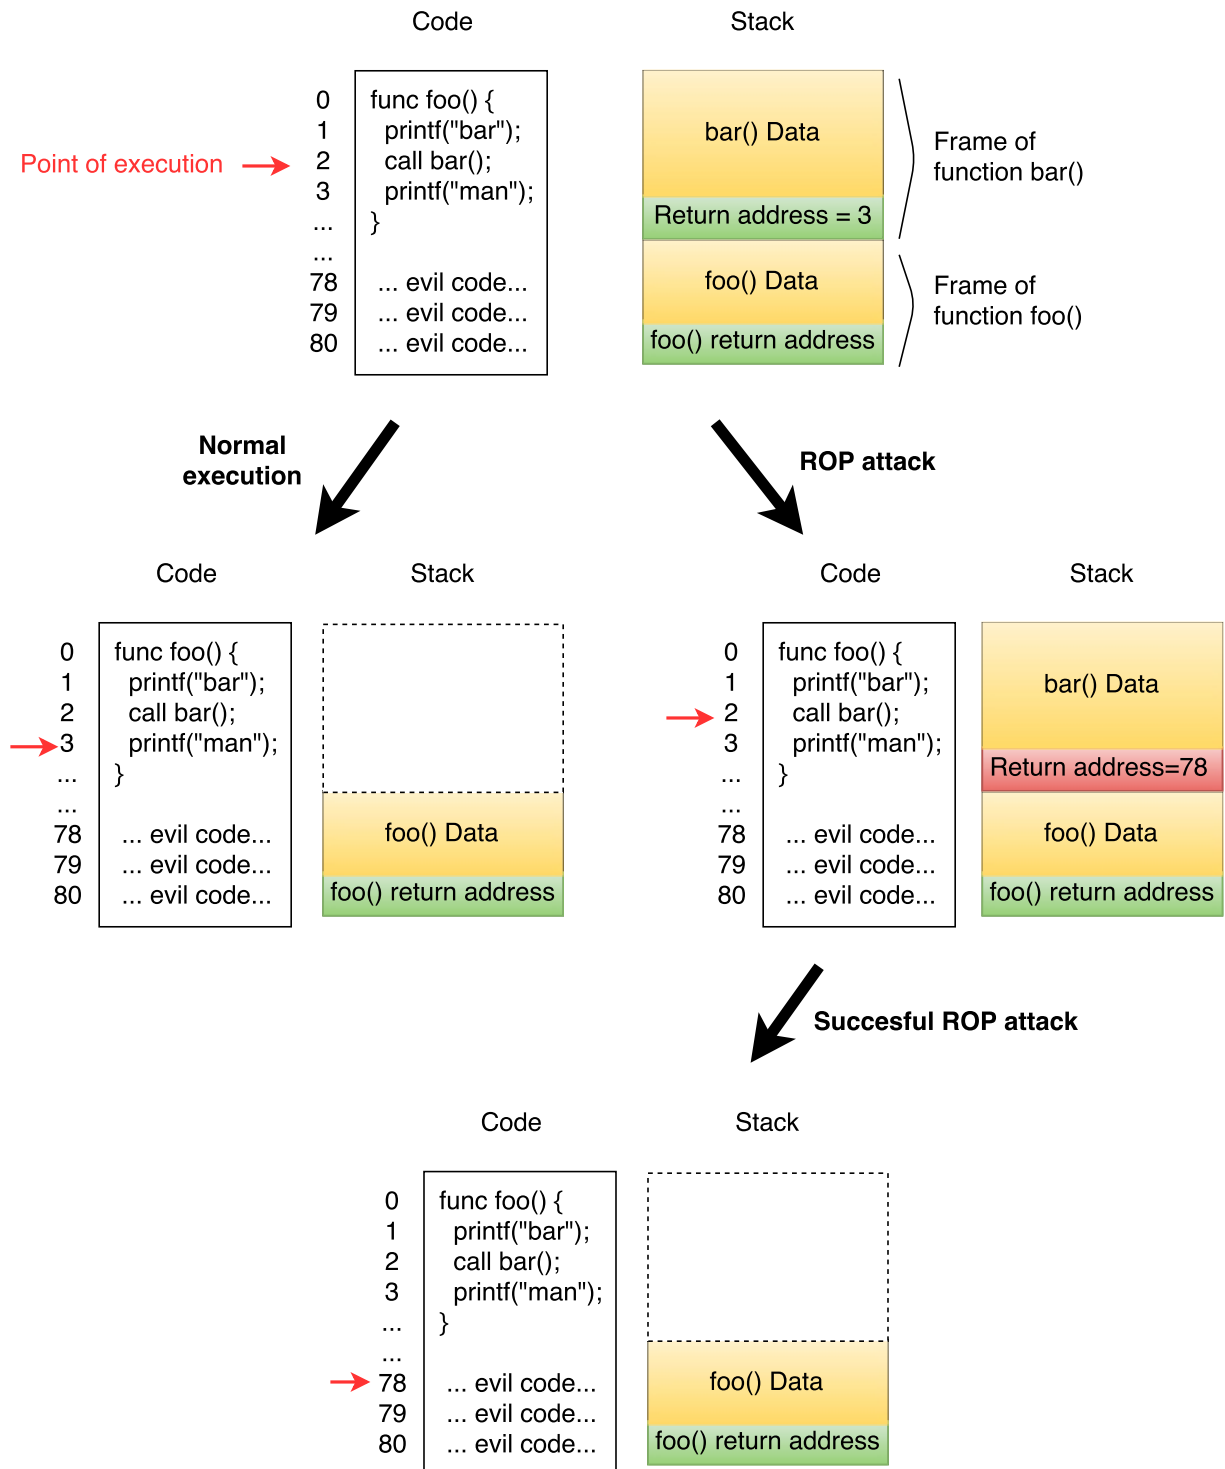


Figure 1: ROP attack




Figure 2: Stack modifications

```

int foo(int a) {
    int* pointer = &a;
    //In pointer we get the address of the parameter a which is in
    //the stack

    *(pointer+4) = 0xff000000;
    //We write the value 0xff000000 in the stack
    //if this location is a return address, the ROP attack succeeds
}

```

Figure 3: Example of dangerous C statement

The second step is to detect every possibly harmful to return addresses (Figure 3). All the operations which can write in the memory can be considered dangerous. Hence we target all instruction related to variable assignment. If we want more precision, we can also target assignments to variables located on the stack.

Finally when we have detected all the dangerous statements we transform the module code. Before each of this dangerous statement we add a protection mechanism similar to masking in SFI. If the address written does not match the location of a return address, then the operation is allowed and the program continues to run casually. Whereas if there is an attempt to write on a return address we trigger an error behaviour like crashing (Figure 4).

To sum it up, our approach aims to have an easy way to know return addresses location and then add a check at runtime before every dangerous instruction to prevent illegal writing on return addresses location. To do this we divided the approach into four phases:

1. Fix stack frames size
2. Align the stack
3. Detect dangerous statements
4. Secure the dangerous statements

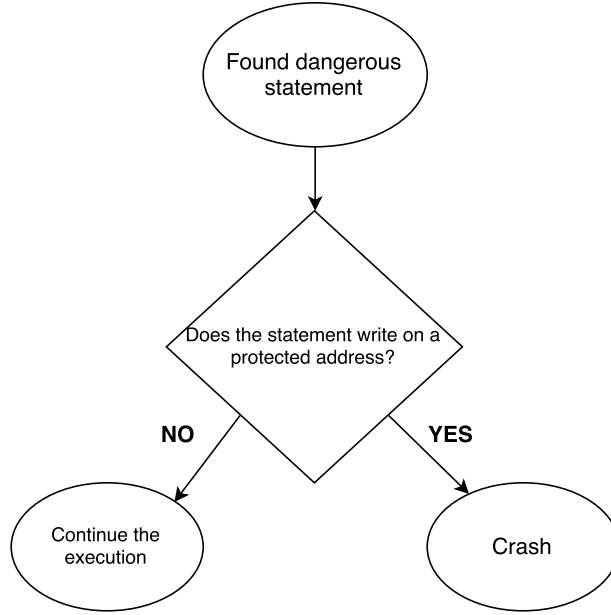


Figure 4: Runtime check algorithm

### 3.2 Security properties

The approach we propose is composed of four phases, to get the confidence that our idea is effective in protecting return addresses we are going to formalize the properties we expect from each phase. Furthermore like we pointed earlier, we're going to work with the certified compiler CompCert. The ideal way to be sure of our idea would be to prove it with Coq the proof assistant the language used to build CompCert. By working with these tools we hope that one day we will be able to prove some security guarantees brought by our approach.

#### 1. Fixed stack frames size

- Return addresses locations are all separated by a constant offset equals to the size of the frames
- The transformation does not change the semantic of the module

#### 2. Stack alignment

- The first return address location of the stack have its least significant bits equals to 0
- The transformation does not change the semantic of the module

#### 3. Detection of memory write statements

- Every statement of the analysed code that might modify the stack memory state is detected

#### 4. Securing memory write statements

- Statements that does not involve the protected addresses will keep their behaviour



TODO substitute with the right syntax when finished

```
int foo(int a) {
    asm(“\sub 50, \esp”);
    //This line does the operation ESP = ESP - 50
    //This disrupts the stack layout we establish in our
    transformation
    printf(“Hello world!”);
}
```

Figure 5: C inline assembly

- The protection will trigger an error behaviour if we try to write on a protected address

1. and 2. combined give us the guarantee that the least significant bits of all the return addresses location will be equal to  $0 \bmod n$  with  $n$  the size of the frames .

We make it so the protection mechanism prevents any write on addresses located in the **stack memory area** and their least significant bits equal to  $0 \bmod n$ .

If all these properties are fulfilled we are confident that under certain conditions it's not possible for a hazardous module to modify the flow control through the return addresses. The necessary conditions for our approach to work will be discussed right after.

### 3.3 Analysis of the approach

#### 3.3.1 Conditions

We believe that the solution we've just presented can bring very strong security properties against ROP attacks. However for this approach to work we need certain hypothesis to be true. Indeed some of the properties enumerated before become false after certain operations.

- **Stack modifications**, every operation that disrupts the stack structure may nullify our property that says "every return addresses are separated by a fixed offset". For example x86 architecture use the ESP register to keep track of the stack growth. If we fiddle with it we may introduce a shift in the return addresses location. Then our runtime check on  $0 \bmod n$  addresses would not be relevant anymore. An example of such operation in C is inline assembly which allows us to put some assembly code in C code (Figure 5). One simple solution to this flaw would be to forbid any usage of inline assembly.
- **Unsecure libraries**, for our approach to work we need to have all dangerous write statements to contain our runtime checks. Hence all executed code must have been compiled with our transformation. For example, the *glibc* library of C contains multiple insecure functions like *printf*, *strcpy*... Furthermore those flawed functions are common vulnerabilities for *buffer overflows* attacks which are a type of ROP attack. To avoid this issue we would need to rewrite the *glibc* or compile it with our tools.

Those conditions are necessary for our approach to be relevant. Yet we can't guarantee that these conditions are an exhaustive list of the hypothesis needed. They're the conditions we could think of but there might be some more.

### 3.3.2 Discussion

We have presented the principle of our approach in this chapter. Following we mentioned some necessary conditions for our solution to be relevant. In this section we're going to discuss about the pros, cons or remarks about the proposed solution.

The benefits of our transformation is clear, any code compiled with a compiler enforcing our methods is unable to interfere with the control flow of our program through return addresses. Furthermore if we combine our solution with the SFI presented earlier we can have some strong security properties on the execution of dangerous modules with our main program. Alas there are also some disadvantages to our approach that we're going to present here:

- **Architecture dependant**, our solution depends a lot of the stack layout of the program. Indeed fixing the size of the frames requires us to modify the original stack layout. Therefore since the stack layout vary depending of the architecture and compiler you're using, the modifications that have to be done are also different. We can then easily comprehend that we would need a different implementations for every existing stack layout. Moreover since these layouts can be really different it might be very gruesome to implement our solution on certain of them. In the implementation we present after we focus solely on x86-32 architecture with the compiler CompCert.
- **Memory consumption**, since we're fixing the size of the frames instead of having dynamic sizes the memory usage of the stack is bigger. We have the issue of choosing an adequate size for the frames in our solution. The easiest one is to take the maximum frame size of the program as the constant size for all the frames. The downside is that we might have a memory usage explosion from our stack. To put a cost on the impact of our solution on the stack size we would need to make numerous tests. Unfortunately during the span of the internship we were unable to do so but it's one of our objectives for the remaining time.

Despite the cons presented we believe the benefits we gain from this method is worth it. With regard to the negative impacts on stack memory consumption we personally didn't encounter any issue with the different tests we made our implementation go through. The impacts may be visible on especially big programs which we did not test yet. We're going to present in the following section the implementation we made based on the ideas we introduced here. This implementation was made with the compiler CompCert for the x86-32 architecture. We're targeting programs written in C, which explains that all the examples we used were related to the C language.

## 4 Implementation

For the implementation of our idea we chose to work with the compiler CompCert. CompCert already had an implementation of SFI presented earlier. Thus if we could combine our approach with the SFI, any program compiled with CompCert would have strong security guarantees. Furthermore CompCert is written with Coq the proof assistant, we eventually hope that we will be

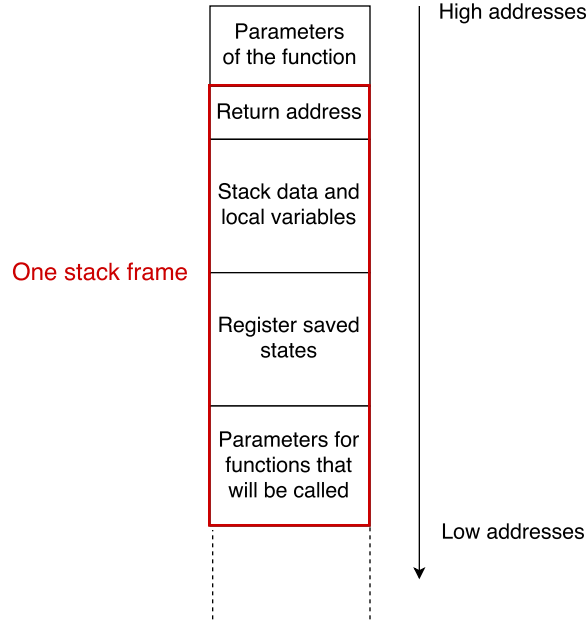


Figure 6: CompCert x86-32 stack layout

able to prove these security properties. In this section we’re going explain in details how we implemented the approach and the different choices we did during the process. Afterwards we will discuss these choices and evaluate the results and performance obtained.

## 4.1 Implementation

Our approach is separated in four phases: “Fixed stack frames size”, “Stack alignment”, “Detection of memory write statements” and “Evaluation of the implementation”. We’re going detail the implementation of these phases in the following sections. These transformations are deeply linked to the stack layout, hence to have a better understanding we’re going to start by introducing the CompCert stack structure.

### 4.1.1 CompCert stack

The layout of the stack is dependent of the architecture and the compiler/interpreter used. For a better understanding of the future sections we describe here the stack layout of x86-32 in CompCert. First of all in x86 the stack grows downwards, it means that the stack grows from the highest addresses to the low ones. A lot of information are saved in the stack, we can find the parameters of the functions, local variables, saved state of registers and the return address of the function. In CompCert the stack layout is composed as in Figure 6.

Each frame is built when a function is called, the different steps related to the creation of a frame is called *function call routine*. CompCert function call routine is described in the Figure 7 and here:

1. Write the return address

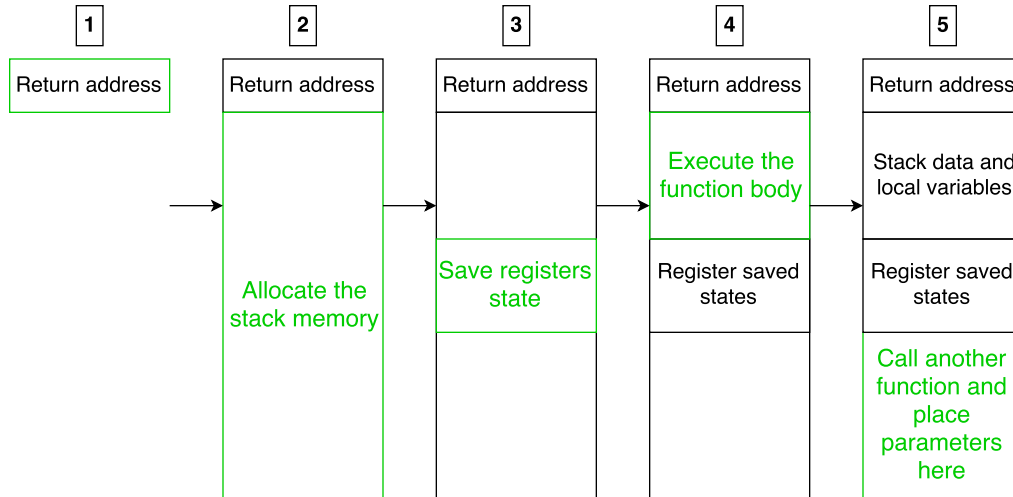


Figure 7: CompCert function call routine

2. Allocate enough memory for the rest of the stack
3. Save registers states in the stack
4. Execute the function body (use the memory for local and stack data)
5. When calling another function, place its parameters at the end of the stack and repeat the process

When returning from a function the routine is pretty much the opposite:

1. Restore registers state
2. Deallocate the stack until the return addresses
3. Pop the return address memory and jump to the value stored in it

#### 4.1.2 Fixed stack frames size

During this phase we want to ensure these two properties:

- Return addresses locations are all separated by a constant offset equals to the size of the frames
- The transformation does not change the semantic of the module

Fortunately in the function call routine of CompCert the return addresses is always the first location in new frames. This peculiarity makes the task easier, indeed, since the location of the return address is fixed in the stack we can deduce that with just fixed size frame the offset between two return addresses locations will be constant. This peculiarity is not always true, for example in x86 architecture with the compiler *gcc* the location of the return addresses changes relatively to the frame depending of the parameters of the function.

To fix the size of stack we had to find the description in of the stack in CompCert. Then we just had to put a constant in the attribute *size* of the stack and readjust the alignment of the different part of a frame. We told CompCert to keep the return address of the stack as the first location in the frame and that all the extra space introduced by the fixed size will be taken for *stack data and locals*. The remaining parts keep the same alignment as before.

#### **4.1.3 Stack alignment**

#### **4.1.4 Detection of memory write statements**

#### **4.1.5 Securing memory write statements**

### **4.2 Evaluation of the implementation**